



Unveiling the Impact of Multi-modal Content in Multi-modal Recommender Systems

Guipeng Xv*

xuguipeng@stu.xmu.edu.cn
School of Informatics, Xiamen
University
Xiamen, China

Xinyu Li*

xinyuli@stu.xmu.edu.cn
School of Informatics, Xiamen
University
Xiamen, China

Yi Liu

yisulphur@stu.xmu.edu.cn
Institute of Artificial Intelligence,
Xiamen University
Xiamen, China

Chen Lin†

chenlin@xmu.edu.cn
School of Informatics, Xiamen
University
Xiamen, China

Xiaoli Wang

xlwang@xmu.edu.cn
Super Intelligent Medical Innovation
Research Center, School of
Informatics, Xiamen University
Xiamen, China

Abstract

Multi-modal recommender systems (MRSs) have emerged as critical multi-modal technologies, but *Do we truly leverage multi-modal content properly?* Through an empirical study of four diverse, real-world datasets spanning various recommendation scenarios, we observe that MRSs exhibit a stronger tendency to recommend items with high similarity to users' past interactions in terms of multi-modal content than conventional RSs. While this tendency improves the recommendation accuracy, it introduces a previously unexplored bias that significantly impacts user experience. We define this bias as **User-side Content Bias**: *users who prefer items similar to their historical choices receive higher quality recommendations than those seeking diverse options*. We show that User-side Content Bias is unrelated to the activity of users, indicating a fundamental limitation in current MRSs. We propose **ISOLATOR**: utIlizing uSer-side cOntent simiLarity via a model-AgnosTic framewORk to leverage multi-modal content more properly. ISOLATOR estimates the impact of User-side Content Similarity and proposes two intervention strategies to meet the needs for more accurate and unbiased recommendations. Extensive evaluations on several widely used datasets demonstrate that ISOLATOR consistently improves various state-of-the-art MRSs and effectively addresses the User-side Content Bias.

CCS Concepts

- Information systems → Multimedia and multimodal retrieval; Recommender systems.

*The authors contribute equally.

†Corresponding author.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

MM '25, Dublin, Ireland

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 979-8-4007-2035-2/2025/10
<https://doi.org/10.1145/3746027.3755300>

Keywords

Multi-modal Recommender System, User-side Content Bias, Causal Inference

ACM Reference Format:

Guipeng Xv, Xinyu Li, Yi Liu, Chen Lin, and Xiaoli Wang. 2025. Unveiling the Impact of Multi-modal Content in Multi-modal Recommender Systems. In *Proceedings of the 33rd ACM International Conference on Multimedia (MM '25)*, October 27–31, 2025, Dublin, Ireland. ACM, New York, NY, USA, 10 pages. <https://doi.org/10.1145/3746027.3755300>

1 Introduction

Multi-modal Recommender Systems (MRSs) have garnered considerable interest in recent years [16, 21, 37, 45]. It is well regarded [36, 41, 46] that multi-modal content enhances the overall understanding of item and improves user preference prediction compared with RSs using only user feedback or single-modal content. Therefore, as critical multi-modal technologies, MRSs are widely employed on online platforms, such as E-commerce [14, 19, 20, 34] and video sharing [1, 2, 30, 39] platforms.

Current MRSs primarily focus on leveraging advanced techniques to integrate multi-modal and behavioral features, with the goal of enhancing recommendation accuracy [16, 27, 32, 35]. However, there is a critical question that remains insufficiently addressed – *Do we truly leverage multi-modal content properly?*

To answer this question, we conduct a cross-domain empirical study on four real-world, widely-used datasets from E-commerce [23] and short-form video [24] platforms. We find that existing MRSs incorporate multi-modal content **with positive and negative impacts**. (1) On the one hand, compared with conventional RSs, MRSs recommend items more similar to users' past interactions regarding multi-modal content. We observe that the overall recommendation performance (i.e., recall rate) positively correlates with the multi-modal content similarity between items in the recommendation list and the user's historical interactions (Sec. 3.3). (2) On the other hand, including multi-modal content threatens the fairness of recommendations from the user's perspective. We create user groups, i.e., **Content-Consistent Users** (who prefer items similar to their past interactions) and **Content-Diverse Users** (who prefer items

distinct from their past interactions). We observe that MRSs amplify the performance gap between Content-Consistent Users and Content-Diverse Users, referred to as User-side Content Bias (Sec. 3.4). This phenomenon brings a new fairness issue, i.e., User-side Content Bias is independent of the revealed unfairness on inactive users [15]. With the same number of interactions, Content-Consistent Users still receive higher quality recommendations than Content-Diverse Users (Sec. 3.5).

Given the widespread presence of MRSs of different model architectures, it is crucial to develop a generally applicable framework to balance the positive and negative impacts of multi-modal content, considering the specific goals and context of the RSs. For example, when prioritizing equal experience for all users, User-side Content Similarity should be discarded completely to avoid User-side Content Bias. On the contrary, when recommendation accuracy is pivotal in user satisfaction and engagement, User-side Content Similarity should be selectively utilized to grasp the item relations and improve overall accuracy.

To properly leverage multi-modal content, we construct a causal graph [26] to analyze how the User-side Content Similarity influences the recommendations, and propose **ISOLATOR**: utIlizing uSer-side cOntent SimiLarity via a Model-AgnosTic FrameWORk. In the training stage, ISOLATOR estimates the impact of User-side Content Similarity on recommendation accuracy and employs a *do-calculus* [26] for each user-item pair. In the inference stage, ISOLATOR proposes two strategies tailored for different recommendation scenarios: a Debiasing intervention strategy designed to *eliminate the impact* for all users to mitigate the User-side Content Bias; and a User-specific intervention strategy, which *leverages the impact* by applying personalized interventions for different users to enhance accuracy.

Our main contributions can be summarized as follows:

- (1) We find that multi-modal content leads the MRSs to recommend items more similar to users' past interactions, resulting in amplifying User-side Content Bias. To our knowledge, **this is the first work to investigate User-side Content Bias in MRSs**.
- (2) We propose ISOLATOR to use multi-modal content properly. This light-weight and model-agnostic framework encompasses two strategies, one for debiasing and the other for performance enhancement, to meet different recommendation requirements. To our knowledge, **this is the first work to mitigate User-side Content Bias in MRSs**.
- (3) Extensive experiments on several real-world datasets and widely-used MRSs verify that ISOLATOR not only improves the recommendation performance but also effectively mitigates the bias.

2 Related Work

Multi-modal Recommender Systems. Multi-Modal Recommender Systems (MRSs) predict user preferences by fusing multi-modal and behavioral features, primarily categorized into two frameworks: (1) Matrix Factorization (MF)-based methods [3, 8], which integrate multi-modal item characteristics into MF training. (2) Graph Collaborative Filtering (GCF)-based methods, which can be further categorized into two subtypes. The first subtype integrates multi-modal features into the user-item graph [32, 35, 36], while the second leverages multi-modal content to construct item-item similarity graphs for embedding refinement [18, 27, 37, 45].

Table 1: Statistics of the datasets. $|\mathcal{D}|$, $|\mathcal{U}|$, and $|\mathcal{I}|$ represent the number of interactions, users, and items. \bar{S}^v and \bar{S}^t represents the average visual and textual similarity.

Datasets	$ \mathcal{D} $	$ \mathcal{U} $	$ \mathcal{I} $	\bar{S}^v	\bar{S}^t	Sparsity
Baby	160,792	19,445	7,050	0.2239	0.2626	99.88%
Sports	296,337	35,598	18,357	0.2183	0.2084	99.95%
Clothing	278,677	39,387	23,033	0.2239	0.3880	99.97%
Microlens	705,174	98,129	17,228	0.5078	0.3822	99.96%

Bias in Multi-modal Recommender Systems. Biases have been widely studied in recommender systems, particularly those arising from feedback data and conventional recommendation algorithm [4, 13, 17, 22, 38]. However, **biases introduced by multi-modal content are more critical in MRSs**, which can be categorized into *inter-modality bias* and *item-side bias*. Inter-modality bias arises from challenges in integrating diverse modalities, such as MRSs overly rely on dominant modalities [21], leading to suboptimal recommendations. Item-side bias, on the other hand, occurs when multi-modal content causes unfair phenomena to items, such as over-recommendation of items with specific content [29] or increasing item-side popularity bias [22].

Remarks. ISOLATOR systematically explores the impact of multi-modal content and discovers a new type of bias, i.e., **User-side Content Bias**. Different from other biases and group fairness [10, 11, 15, 42] in RS, ISOLATOR considers **the bias from multi-modal content in user-side** due to varying user preferences for multi-modal content. To our knowledge, ISOLATOR is **the first work** that explores the impact of multi-modal content on the user side and **mitigates the User-side Content Bias in MRSs**.

3 Empirical Study

3.1 Preliminaries

Let $\mathcal{U} = \{u_1, u_2, \dots, u_{|\mathcal{U}|}\}$ and $\mathcal{I} = \{i_1, i_2, \dots, i_{|\mathcal{I}|}\}$ denote the set of users and items, respectively. $|\mathcal{U}|$ and $|\mathcal{I}|$ is the number of users and items. For each user u , let $\mathcal{I}_u \subseteq \mathcal{I}$ be the set of items that u has interacted with in the training set, and let $|\mathcal{I}_u|$ denote its size. Each item i has multi-modal content features $e_i^m \in \mathbb{R}^{d_m}$, where d is the dimension of the features, $m \in \mathcal{M}$ is the modality¹, and \mathcal{M} is the set of modalities. To facilitate similarity computations, these content feature vectors are typically L_2 -normalized, i.e., $\|e\|_2 = 1$.

The goal of multi-modal recommender systems (MRSs) is to generate a ranked list of potential recommendations that each user u may prefer by predicting the user-item preference score $\widehat{r}_{u,i}$ using both behavior and multi-modal content. Formally,

$$\widehat{\mathcal{I}}_u = \text{Top}@k_{i \in \{\mathcal{I} \setminus \mathcal{I}_u\}} \widehat{r}_{u,i}, \quad (1)$$

where $\widehat{\mathcal{I}}_u$ denotes the recommendation list for user u , and the length of recommendation list $|\widehat{\mathcal{I}}_u| = k$.

3.2 Empirical Study Protocol

3.2.1 Datasets. We conduct experiments on four widely-used public datasets following [16, 18, 27, 37, 40, 45], including three Amazon collections (namely Baby, Sports, and Clothing) [23] and the

¹We use textual and visual modalities in this work, but it can also be extended to other modalities.

Microlens [24] dataset. These datasets were chosen due to several advantageous characteristics: (1) First, they undergo a public unified preprocessing pipeline that *standardizes both the dataset structure and the modal content*²; (2) Second, they originate from E-commerce platforms and short-form video platforms, which are two representative multi-modal recommendation scenarios; (3) Third, they offer diverse characteristics in terms of size, sparsity, and multi-modal properties. These advantages enable **more robust and fair comparisons**. The datasets are filtered using a 5-core criterion to ensure sufficient feedback for each user and item and are split into training, validation, and test sets in an 8:1:1 ratio. We compute the average pairwise cosine similarity across all items in each modality, denoted as \bar{S}^v for images and \bar{S}^t for text. Table 1 lists the statistics of the four datasets.

3.2.2 Evaluation metrics. We employ *Recall@20* to measure the recommendation accuracy [8, 18, 27, 32, 35]. *Recall@20* computes the ratio of relevant items for a user in the testing set appearing in the top twenty recommendations made by the baseline. Higher values of *Recall@20* indicate more accurate recommendation results³.

3.2.3 Baselines. Multi-modal recommender systems (MRSs) typically integrate multi-modal content into conventional RSs based on two primary backbone architectures: matrix factorization (MF) and graph collaborative filtering (GCF). We use two conventional RSs MFBPR [28] and LightGCN [9], and five MRSs VBPR [8], GRCN [35], MGNC [27], SLMRec [32] and GUME [18] as backbone since they are widely used. It should be noted that *MFBPR* and *VBPR* are *MF architectures*, while the remaining are *GCF architectures*. To ensure optimal performance, we use publicly available code⁴ [44] with default parameter settings. We train the models using the training set, save the best model based on *Recall@20* on the validation set, and use the models to deliver recommendations on the testing set.

3.3 Positive Impact of Content Similarity

User-side Content Similarity. Essentially, multi-modal recommender systems (MRSs) match items with a user by taking into account the multi-modal content similarity between the candidate item and the items that the user previously interacted with, which we refer to as User-side Content Similarity.

Thus, to investigate the impact of multi-modal content, we monitor the relation between User-side Content Similarity and recommendation performance on different baselines. We use the average content similarity of the recommendation list to users' past interactions (referred to as Average Content Similarity ρ) to reflect the overall recommendations of the model.

$$\rho = \frac{\sum_{u \in \mathcal{U}} \left(\sum_{m \in \mathcal{M}} \left(\sum_{i \in \mathcal{I}_u} \sum_{j \in \widehat{\mathcal{I}}_u} (\text{sim}(\mathbf{e}_i^m, \mathbf{e}_j^m)) \right) \right)}{|\mathcal{U}| \times |\mathcal{M}| \times |\mathcal{I}_u| \times |\widehat{\mathcal{I}}_u|}, \quad (2)$$

where $\text{sim}(\cdot)$ is the pairwise cosine similarity, \mathbf{e}^m is the raw pre-extracted fixed features of modality m , which is model-agnostic. ρ is averaged over all users, all modalities, and all items in the recommendation list for a user. A higher ρ signifies that the recommended items are closer to the user's past interactions regarding

²<https://github.com/enoche/MMRec/tree/master/data>

³Similar trends are also observed when using *NDCG*.

⁴<https://github.com/enoche/MMRec>

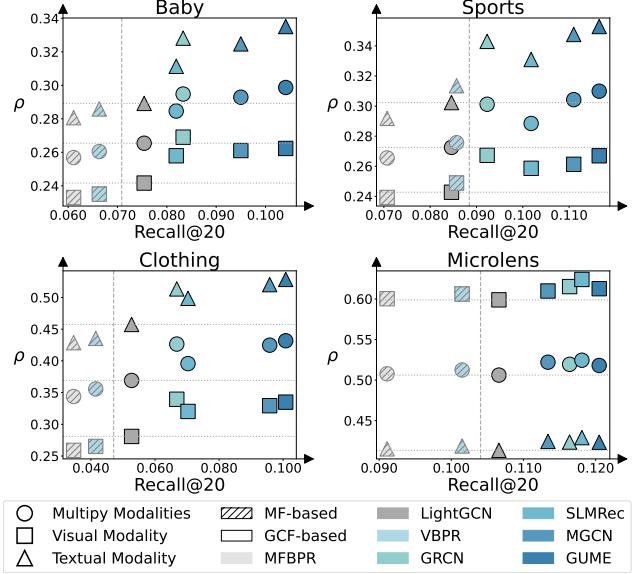


Figure 1: Relationship Between *Recall@20* and the average content similarity of recommendation lists to users' past interactions ρ in different modalities.

multi-modal content, meaning that the baseline is more reliant on multi-modal content similarity.

In Fig. 1, the blue markers represent multi-modal recommender systems (MRSs) while the gray markers represent conventional recommender systems. We have the following observations:

- Comparing methods based on the same backbone architecture, MRSs tend to recommend items with higher content similarity than conventional RSs. For example, the blue markers (MRSs) are typically positioned higher than the gray markers (conventional RSs) of the same shape in Fig. 1, showing that MRSs focus on utilizing content similarity more. This observation is consistent across all tested baselines and datasets, and the similarity can be measured based on either single or multiple modalities.
- MRSs consistently exhibit higher recommendation accuracy than conventional RSs using the same backbone architecture. The blue markers (MRSs) are generally positioned to the right of the gray markers (conventional RSs) of the same shape in Fig. 1, showing that MRSs outperform conventional RSs regarding the *Recall@20* across all datasets.
- Higher values of Average Content Similarity generally lead to improved recommendation accuracy. In the Baby, Sports, and Clothes datasets, there is a positive correlation between accuracy (*Recall@20*) and Average Content Similarity (ρ). In the Microlens dataset, the largest ρ is observed on the best-performing baselines (i.e., MGNC and GUME) and the smallest ρ on the worst-performing baseline (i.e., MFBPR).
- Multi-modal content can more stably reflect content similarity than single-modal content. Some datasets show greater similarities in textual contents, while in other datasets, the image contents are more similar. Specifically, textual similarity (triangular markers in Fig. 1) is relatively high in e-commerce scenarios but lower in short-form video scenarios. Therefore, using single-modal content similarity is inaccurate for certain datasets, and

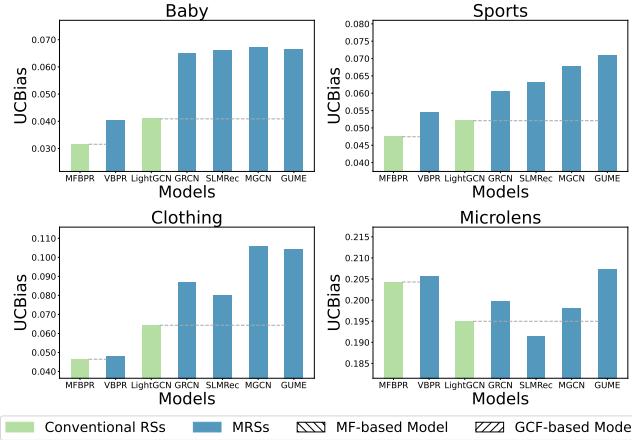


Figure 2: User-side Content Bias between Content-Consistent Users and Content-Diverse Users.

we should use multi-modal content similarity as a domain-robust assessment.

3.4 Negative Impact of Content Similarity

We have shown that the overall recommendation performance is improved by recommending items that have a higher multi-modal content similarity with the user’s past interactions. However, users show significantly different preference patterns. Some users prefer items with highly similar content, such as those who favor consistent vintage styles in clothing, accessories, and home goods. In contrast, some users prefer exploring items with diverse content, such as those who enjoy transcending different styles and combining various items that may seem unrelated at first glance. In this subsection, we investigate the negative impact of User-side Content Similarity, i.e., whether it will cause biased performance on different users.

User Grouping. We divide users according to their preference patterns. Firstly, we compute the average pairwise content similarity of interacted items for each user (denoted as c_u):

$$c_u = \frac{\sum_{m \in \mathcal{M}} \sum_{i \in I_u} \sum_{j \in I_u, i \neq j} (\text{sim}(\mathbf{e}_i^m, \mathbf{e}_j^m))}{|\mathcal{M}| \times |I_u| \times (|I_u| - 1)}, \quad (3)$$

where $\text{sim}(\cdot)$ is the pairwise cosine similarity.

Secondly, based on c_u , we identify two user groups:

- **Content-Consistent Users:** the top $N\%$ users with the greatest c_u . Their past interactions show greater consistency in content, indicating a more focused user interest.
- **Content-Diverse Users:** the bottom $N\%$ users with the lowest c_u . Their past interactions show greater diversity in content, indicating a more varied user interest.

We use $N = 20$ in this paper, but we consistently observe a similar situation as N varies across different coverage ranges, extending from 5% to 25%.

User-side Content Bias. Following [6, 17], we propose a metric to quantify the biased treatment for different user groups: User-side Content Bias (denoted as $UCBias$) is defined as the gap between recommendation performance on Content-Consistent Users and Content-Diverse Users, which is calculated as

$$UCBias = E(U^C) - E(U^D), \quad (4)$$

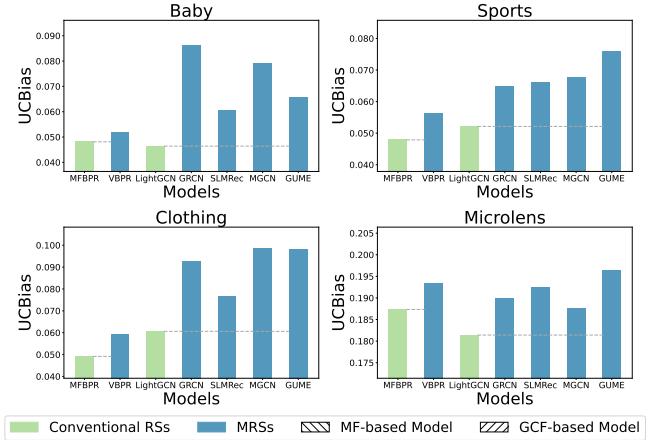


Figure 3: User-side Content Bias between Content-Consistent Users and Content-Diverse Users with the same number of interactions.

where $E(\cdot)$ is denoted as the evaluation metrics. We use $Recall@20$ for $E(\cdot)$ in this paper. U^C and U^D are the Content-Consistent Users and Content-Diverse Users groups, respectively.

We can observe from Fig. 2 that

- Regardless of the type of recommendation models, whether they are shallow or deep, multi-modal or non-multi-modal, and irrespective of the context or distribution of the datasets, **there exists a positive performance gap, indicating that Content-Consistent Users receive higher quality recommendations than Content-Diverse Users.**
- When comparing methods within the same backbone architecture, the blue bars (represent MRSs) are typically higher than the green bars (represent conventional RSSs), indicating that multi-modal content amplifies the performance gap.

User-side Content Bias brings a serious unfairness issue. It makes users see more items similar to their past interactions. For Content-Consistent Users, this means better matches and a great experience. But for Content-Diverse Users, it limits diversity, worsens the filter bubble [31], and results in poor recommendations. This unfairness can drive Content-Diverse Users away, harming the overall performance and reputation of the recommendation systems. Furthermore, since Content-Diverse Users has the same size as Content-Consistent Users in our study, the User-side Content Bias affects a significant amount of users.

3.5 Uniqueness of User-side Content Bias

Our next question is whether User-side Content Bias is a new type of bias. Since User-side Content Bias is evaluated from the user’s perspective, we need to distinguish it from the ϵ -fairness, i.e., inactive users receive unfair treatment [15].

To show that User-side Content Bias is unrelated to ϵ -fairness [15], we first compute the median number of interactions per user in each dataset (denoted as a) to avoid skewing the analysis. Then, we randomly sample the same number of users who have interacted with exactly a items from both Content-Consistent Users and Content-Diverse Users groups. We report the results within these subgroups.

Algorithm 1: ISOLATOR

Input: Dataset \mathcal{D} ; Hyper-parameter α, β, γ, t ; Inference intervention strategy: *D-ISOLATOR* or *U-ISOLATOR*;
Output: Recommendation list;

// Training Stage.

- 1 Calculate User-side Content Similarity $s_{u,i}$ by Equ. 9;
- 2 Estimate the probability of the interaction by Equ. 7;
- 3 Update parameters Θ by optimizing Equ. 5;

// Inference Stage.

- 4 **if** D-ISOLATOR **then**
- 5 Estimate the probability of the interaction by Equ. 10;
- 6 Generate recommendation lists by Equ. 1;
- 7 **else**
- 8 Estimate the impact of User-side Content Similarity that users require $\tilde{s}_{u,i}$ by Equ. 11;
- 9 Estimate the probability of the interaction by Equ. 13;
- 10 Generate recommendation lists by Equ. 1;

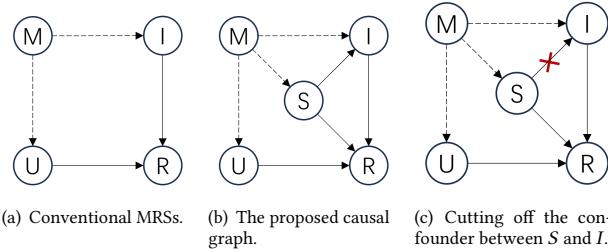


Figure 4: Causal graphs to describe the recommendation process. U: user, I: item, M: multi-modal content, R: interaction probability, S: User-side Content Similarity.

From Fig. 3, we can observe that User-side Content Bias consistently exists regardless of the dataset used or the method applied. Moreover, MRSs consistently exhibit a larger User-side Content Bias. The results support that User-side Content Bias is independent of the number of users' past interactions.

4 Methodology

To effectively leverage the multi-modal content in MRSs, we first analyze the impact of User-side Content Similarity on multi-modal recommender systems (MRSs) from a causal view. Then, we propose **ISOLATOR**: utilizing uSer-side cOntent simiLaritY via a model-AgnosTic framewORk to achieve more accurate and unbiased recommendations. ISOLATOR estimates and disentangles the impact of User-side Content Similarity in the training stage and employs two innovative intervention strategies in the inference stage to address its side effects. These strategies focus on debiasing and performance enhancement, respectively, to cater to various recommendation requirements. The learning program of ISOLATOR is summarized in Algorithm 1.

4.1 Causal View of MRSs

4.1.1 Causal View of Conventional MRSs. As shown in Fig. 4(a), conventional multi-modal recommender systems (MRSs) typically use four variables: user (U), exposed item (I , including exposure and features attributes), multi-modal content (M), and interaction label (R). $R = 1$ represents that the user had interacted with the item. The edges represent the causal relations between variables.

- Edges $M \rightarrow I$ and $M \rightarrow U$ signify that multi-modal content forms the feature of items and users.
- Edge $\{U, I\} \rightarrow R$ denotes that matching between user preferences U and item attributes I dictates user-item interaction R .

Based on above insights, conventional MRSs model the probability function $P(R|U, I)$ ⁵ by the parameterized function $f_\Theta(u, i)$, which means given a user-item pair $U = u, I = i$, how likely the user u will interact with the item i . Specifically, they optimize the model parameters Θ of RSs on dataset \mathcal{D} via the widely-used Bayesian Personalized Ranking [28] (BPR) loss during the training stage:

$$\mathcal{L}_{BPR} = \sum_{(u, i, j) \in \mathcal{D}} \ln \sigma(P_\Theta(r=1|u, i) - P_\Theta(r=1|u, j)) + \lambda_\Theta \|\Theta\|_2, \quad (5)$$

where i denotes the positive sample for user u , j denotes the negative sample for user u , $\sigma(\cdot)$ is the sigmoid function, the $\|\cdot\|_2$ is the L2-regularization, and λ_Θ is the regularization coefficient.

4.1.2 Causal View of ISOLATOR. Unlike conventional MRSs, our method incorporates a new latent variable User-side Content Similarity (S) in Fig. 4(b), which includes new causal relations between variables.

- Edge $M \rightarrow S$ denotes that multi-modal content can calculate the User-side Content Similarity.
- Edge $\{U, I, S\} \rightarrow R$ denotes that an interaction label R is influenced by the user U , the item I , and the User-side Content Similarity S . We introduce a cause node S to capture users' preferences for User-side Content Similarity, which is because many users interact with items solely based on User-side Content Similarity. Thus, the occurrence of interaction is a synergistic result of U, I , and S .
- Edge $S \rightarrow I$ denotes that User-side Content Similarity influences the exposure of the items. For example, we show that MRSs recommend items with greater User-side Content Similarity to users' past interactions in Sec. 3.3.

According to the causal theory, since User-side Content Similarity S affects both item I and observed interactions R , it acts as a confounder [26], leading to spurious associations if not properly controlled [26]. There are two causal paths: $S \rightarrow R$, and $S \rightarrow I \rightarrow R$. The first path clearly shows such similarity directly impacts the interaction probabilities. The second path shows such similarity indirectly impacts the interaction probabilities by impacting the exposure likelihood of items, making the recommendations more similar to the user's past interactions. However, this effect can amplify the User-side Content Bias. As discussed in Sec. 3.4, not all users are interested in items similar to their past preferences. This effect caters better to the preferences of users who seek similar items (Content-Consistent Users) but neglects the needs of users who seek diverse items (Content-Diverse Users). Consequently, it exacerbates the performance gap between Content-Consistent Users and Content-Diverse Users in MRSs, resulting in unfair recommendations. In conclusion, the path $S \rightarrow I$ has adverse effects and should be removed from MRSs.

To properly leverage multi-modal content and achieve unbiased and high-quality recommendations, it is critical to accurately estimate the causal effect of similarity (User-side Content Similarity)

⁵For simplicity, we will not explicitly emphasize the paths $M \rightarrow I$, $M \rightarrow U$ and $M \rightarrow S$ in our subsequent formulas.

while controlling for potential confounding effect. Inspired by [33, 43], ISOLATOR leverages *do-calculus* [26] to estimate the causal effect of User-side Content Similarity directly from observational data. *Do-calculus* leverage the causal graph to **adjust for confounding variables** and simulate interventions by **fixing certain variables** (e.g., setting User-side Content Similarity) while allowing others to vary naturally, estimating the causal effect. ISOLATOR employs the *do-calculus* to cut off the path $S \rightarrow I$ in training stage, while intervening the path $S \rightarrow R$ in inference stage.

4.2 Training Stage

ISOLATOR performs a *do-calculus* [26] to eliminate the confounding effect of User-side Content Similarity S on item attributes I . This can be understood as cutting off the edge $S \rightarrow I$ in the causal graph, thereby blocking the influence of S on I as shown in Fig. 4(c). We formulate the predictive model as

$$\begin{aligned} P(R|do(U, I)) &\stackrel{(1)}{=} P(R|U, do(I)) \\ &\stackrel{(2)}{=} \sum_s P(R|U, do(I), s)P(s|U, do(I)) \\ &\stackrel{(3)}{=} \sum_s P(R|U, I, s)P(s), \end{aligned} \quad (6)$$

where Equ. 6 (1) is because there is no path between U and R in our causal graph that requires intervention.; Equ. 6 (2) is because of Bayes' theorem; Equ. 6 (3) is because that $do(I)$ cuts off $S \rightarrow I$, so U and I are independent with S .

Based on the formula above, we first estimate $P(R|U, I, S)$. We focus on disentangling the user-item matching and User-side Content Similarity in the training stage, which simplifies the application across MRSs and avoids re-evaluate the entire model in the inference stage. We ensure that the model outputs are monotonically increasing without restricting them to $[0, 1]$, which reduces unnecessary computational costs while avoiding impact item rankings. We design it as:

$$P_\Theta(R|U = u, I = i, S = s_{u,i}) = ELU'(f_\Theta(u, i)) \times (\sigma(s_{u,i}))^\gamma, \quad (7)$$

where $f_\Theta(u, i)$ represents any MRSs. Following [43], we use a variant of the Exponential Linear Unit [5] activation function $ELU'(\cdot)$ to ensure the matching score is positive. The bigger $f_\Theta(u, i)$ is, the bigger $ELU'(f_\Theta(u, i))$ is:

$$ELU'(x) = \begin{cases} e^x, & \text{if } x \leq 0 \\ x + 1, & \text{else} \end{cases} \quad (8)$$

Hyper-parameter $\gamma > 0$ is used to adjust the User-side Content Similarity $s_{u,i}$, a larger value increases its impact. The User-side Content Similarity $s_{u,i}$ is calculated by

$$s_{u,i} = \frac{\sum_{m \in \mathcal{M}} \sum_{j \in \widehat{\mathcal{I}_u}} (\text{sim}(\mathbf{e}^m, \mathbf{e}_j^m))}{|\mathcal{M}| \times |\widehat{\mathcal{I}_u}|}, \quad (9)$$

where \mathbf{e} is multi-modal content features that remain unchanged during training. Therefore, $s_{u,i}$ can be precomputed and stored before training to reduce computational cost. $\sigma(\cdot)$ is the sigmoid activation function to ensure such content similarity is positive. The bigger $\sigma(s)$ is, the bigger $(\sigma(s))^\gamma$ is.

Since the $s_{u,i}$ is an unique constant for each user-item pair, we can estimate $P(R|do(U = u, I = i))$ by $P_\Theta(R|U = u, I = i, S = s_{u,i})$. ISOLATOR also employs the Bayesian Personalized Ranking [28] loss in Equ. 5 to train the model.

4.3 Inference Stage

After disentangling the impact of User-side Content Similarity in the training stage, we eliminate such impact by doing the intervention $P(R|do(U, I), do(S))$ in the inference stage. We set the value of User-side Content Similarity S while allowing others to vary naturally [26]. We propose two strategies to meet different recommendation scenarios.

4.3.1 Debiasing intervention strategy. When we strive to **reduce bias as much as possible while maintaining performance**, we can adopt a Debiasing intervention strategy (short for *D-ISOLATOR*). This strategy applies the same intervention to all users to completely eliminate the impact of User-side Content Similarity. The formula of the Debiasing intervention strategy can be expressed as:

$$P(R|do(U, I), do(S)) = ELU'(f_\Theta(u, i)). \quad (10)$$

Since we disentangle user-item matching and User-side Content Similarity in the training stage, we can efficiently remove such impact without re-training the entire model.

4.3.2 User-specific intervention strategy. The *D-ISOLATOR* eliminates the impact of User-side Content Similarity and applies the same intervention to all users. However, as discussed in Sec. 3.4, users have significantly different preferences for multi-modal content. Thus, when we pursue the goal of **enhancing performance without significantly increasing bias**, we can adopt an User-specific intervention strategy (short for *U-ISOLATOR*) to personalize content similarity for each user.

Since each user has a different preference for multi-modal content, we use the difference between user u 's average pairwise content similarity of past interactions c_u and the User-side Content Similarity $s_{u,i}$ to estimate the impact of User-side Content Similarity that users require.

$$\widehat{s_{u,i}} = g(s_{u,i} - c_u). \quad (11)$$

Specifically, we use the average pairwise similarity of past interactions c_u to reflect the user's expected User-side Content Similarity. There are two different scenarios: (1) When $s_{u,i} - c_u > 0$, it indicates that the current item is similar to past interactions. The smaller the difference, the more closely the current item meets the user's expected User-side Content Similarity. (2) When $s_{u,i} - c_u < 0$, it indicates the current item is different from the past interactions. The larger the difference, the more the current item deviates from the user's expected User-side Content Similarity. These two scenarios should be treated separately because *a positive and negative difference, even with the same magnitude, have different impacts on recommendation accuracy*. Therefore, we construct a generating function $g(\cdot)$ as follows:

$$g(x) = \begin{cases} e^{-\alpha x}, & \text{if } x \geq 0 \\ \frac{4}{1+e^{-\beta x}} - 1, & \text{else} \end{cases} \quad (12)$$

where $\alpha > 0$ and $\beta > 0$ are two hyper-parameters to control the impact; the larger the α , the larger the β , the smaller the impact on the user. When $x = 0$, $g(x)$ reaches its maximum value, which is $g(0) = 1$.

Finally, the formula of the User-specific intervention strategy can be expressed as:

$$P(R|do(U, I), do(S)) = ELU'(f_\Theta(u, i)) \times (\sigma(\widehat{s_{u,i}}))^\iota, \quad (13)$$

where $\iota > 0$ is a hyper-parameter to adjust the $\widehat{s_{u,i}}$.

Table 2: Performance of MRSs with and without ISOLATOR. The best and second results are marked with **Bold** and Underline. * indicates that the p-value is less than 0.05.

Datasets	Baby				Sports				Clothing				Microlens			
	R@10	R@20	N@10	N@20												
MFBPR	0.0388	0.0611	0.0213	0.0271	0.0475	0.0714	0.0256	0.0317	0.0245	0.0353	0.0139	0.0166	0.0591	0.0911	0.0310	0.0393
LightGCN	0.0476	0.0751	0.0256	0.0327	0.0556	0.0848	0.0307	0.0382	0.0353	0.0535	0.0193	0.0240	0.0704	0.1066	0.0367	0.0461
VBPR	0.0414	0.0654	0.0219	0.0282	0.0543	0.0837	0.0295	0.0371	0.0296	0.0441	0.0164	0.0201	0.0660	0.0992	0.0344	0.0430
+D-ISOLATOR	<u>0.0424*</u>	<u>0.0682*</u>	<u>0.0226*</u>	<u>0.0292*</u>	<u>0.0556*</u>	<u>0.0847*</u>	<u>0.0302*</u>	<u>0.0377*</u>	<u>0.0305*</u>	<u>0.0468*</u>	<u>0.0170*</u>	<u>0.0208*</u>	<u>0.0674*</u>	<u>0.1024*</u>	<u>0.0352*</u>	<u>0.0440*</u>
+U-ISOLATOR	0.0453*	0.0715*	0.0243*	0.0311*	0.0559*	0.0858*	0.0305*	0.0382*	0.0322*	0.0485*	0.0175*	0.0216*	0.0684*	0.1038*	0.0356*	0.0448*
GRCN	0.0526	0.0827	0.0284	0.0362	0.0582	0.0890	0.0313	0.0393	0.0431	0.0657	0.0228	0.0286	0.0765	0.1160	0.0399	0.0501
+D-ISOLATOR	<u>0.0531*</u>	<u>0.0839*</u>	<u>0.0287*</u>	<u>0.0367*</u>	0.0605*	0.0925*	0.0328*	0.0411*	<u>0.0440*</u>	<u>0.0665*</u>	<u>0.0233*</u>	<u>0.0289*</u>	<u>0.0772*</u>	<u>0.1168*</u>	0.0403*	0.0505*
+U-ISOLATOR	0.0536*	0.0851*	0.0289*	0.0369*	<u>0.0602*</u>	<u>0.0924*</u>	<u>0.0325*</u>	<u>0.0408*</u>	<u>0.0442*</u>	<u>0.0673*</u>	<u>0.0235*</u>	<u>0.0294*</u>	<u>0.0773*</u>	<u>0.1174*</u>	0.0403*	0.0505*
SLMRec	0.0525	0.0799	0.0285	0.0356	0.0671	0.1010	0.0368	0.0456	0.0461	0.0691	0.0251	0.0309	0.0784	0.1189	0.0406	0.0510
+D-ISOLATOR	<u>0.0545*</u>	<u>0.0851*</u>	<u>0.0295*</u>	<u>0.0373*</u>	<u>0.0674*</u>	<u>0.1013*</u>	<u>0.0371*</u>	<u>0.0459*</u>	<u>0.0466*</u>	<u>0.0696*</u>	<u>0.0253*</u>	<u>0.0312*</u>	<u>0.0786*</u>	<u>0.1193*</u>	<u>0.0407*</u>	<u>0.0511*</u>
+U-ISOLATOR	0.0548*	0.0852*	0.0296*	0.0374*	0.0680*	0.1020*	0.0373*	0.0460*	<u>0.0473*</u>	<u>0.0706*</u>	<u>0.0257*</u>	<u>0.0316*</u>	<u>0.0793*</u>	<u>0.1199*</u>	0.0411*	0.0515*
MGCN	0.0619	0.0958	0.0334	0.0421	0.0744	0.1122	0.0409	0.0506	0.0656	0.0955	0.0361	0.0437	0.0752	0.1138	0.0387	0.0486
+D-ISOLATOR	<u>0.0631*</u>	<u>0.0969*</u>	<u>0.0339*</u>	<u>0.0426*</u>	<u>0.0748*</u>	<u>0.1128*</u>	<u>0.0412*</u>	<u>0.0510*</u>	<u>0.0653</u>	<u>0.0959*</u>	<u>0.0359</u>	<u>0.0438*</u>	<u>0.0756*</u>	<u>0.1143*</u>	<u>0.0390*</u>	<u>0.0489*</u>
+U-ISOLATOR	0.0633*	0.0980*	0.0342*	0.0431*	0.0752*	0.1134*	0.0412*	0.0511*	<u>0.0662*</u>	<u>0.0967*</u>	<u>0.0364*</u>	<u>0.0441*</u>	<u>0.0759*</u>	<u>0.1150*</u>	0.0391*	0.0492*
GUME	0.0682	0.1041	<u>0.0368</u>	0.0460	0.0776	0.1168	0.0425	<u>0.0526</u>	0.0688	0.1008	0.0373	0.0454	0.0813	0.1204	0.0426	<u>0.0526</u>
+D-ISOLATOR	<u>0.0690*</u>	<u>0.1062*</u>	<u>0.0368</u>	<u>0.0464*</u>	<u>0.0778*</u>	<u>0.1171*</u>	<u>0.0427*</u>	<u>0.0528*</u>	<u>0.0690*</u>	<u>0.1010*</u>	<u>0.0374*</u>	<u>0.0455*</u>	<u>0.0817*</u>	<u>0.1207*</u>	0.0428*	0.0528*
+U-ISOLATOR	0.0695*	0.1063*	0.0371*	0.0465*	0.0779*	0.1173*	0.0427*	0.0528*	<u>0.0694*</u>	<u>0.1019*</u>	<u>0.0376*</u>	<u>0.0458*</u>	<u>0.0819*</u>	<u>0.1211*</u>	<u>0.0428*</u>	<u>0.0528*</u>

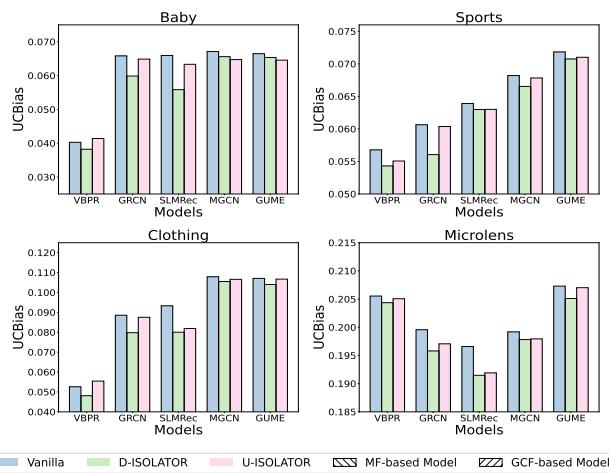


Figure 5: User-side Content Bias in MRSs with and without ISOLATOR. Lower UCBias for better de-biasing effectiveness.

5 Experiments

We use the datasets described in Sec. 3.2. We implement our method in PyTorch [25]. The embedding dimension d is fixed to 64 for all models to ensure a fair comparison. We optimize all models with the Adam [12] optimizer, where the batch size is fixed at 2,048. We use the Xavier initializer [7] to initialize the model parameters. We train the model using the same learning rate as the backbone model. The optimal hyper-parameters are determined via grid search on the validation set: the γ in Equation 7 and the ι in Equation 13 are tuned amongst {0.001, 0.01, 0.1, 1}, the α in Equation 12 is tuned amongst {2, 4, 6, 8, 10}, the β in Equation 12 is tuned amongst {1, 5, 10, 15, 20}. For convergence consideration, the early stopping and total epochs are fixed at 25 and 1,000, respectively. **We make our code available online to ease reproducibility**⁶.

⁶<https://github.com/XMUDM/ISOLATOR>

5.1 Impacts on Recommendation Accuracy

To explore the impact of ISOLATOR on recommendation accuracy, we apply ISOLATOR to various multi-modal recommender systems (MRSs) mentioned in Sec. 3.2.3. We use $Recall@K$ and $NDCG@K$ (abbreviated as R and N) as evaluation metrics following prior works [8, 18, 27, 32, 35]. Higher values of $Recall$ and $NDCG$ indicate more accurate recommendation results. Here we set $K = 10$ and $K = 20$. We have following observations from Table 2:

- ISOLATOR can be applied to various MRSs and consistently improve the recommendation accuracy. We observe consistent improvements across all backbone models in terms of $Recall$ and $NDCG$ on all datasets. This highlights its ability to **effectively capture user preferences for multi-modal content by intervening in the impact of User-side Content Similarity**.
- *U-ISOLATOR* improves recommendation performance by leveraging User-side Content Similarity based on individual user needs, outperforming *D-ISOLATOR* in most cases with superior $Recall$ and $NDCG$. This demonstrates that **leveraging the impact of User-side Content Similarity properly can boost recommendation accuracy**.
- The performance of MRSs typically surpasses that of conventional RSs using the same backbone architecture. This indicates that **incorporating multi-modal content to comprehend the similarity between items can enhance accuracy**.

5.2 Impacts on User-side Content Bias

We investigate the effects of ISOLATOR on mitigating User-side Content Bias in this subsection. We use the top and bottom 20% users as Content-Consistent Users and Content-Diverse Users following Sec. 3.4. We use the $Recall@20$ to calculate the UCBias. A lower UCBias suggests a smaller bias.

As shown in Fig. 5, we have following observations:

- *D-ISOLATOR* effectively mitigates User-side Content Bias, which consistently achieves the lowest UCBias across all models and datasets, reducing the bias by up to 15.33% compared

Table 3: Performance and User-side Content Bias of GUME with and without ISOLATOR in each modality on Baby dataset. The best and second results are marked with Bold and Underline.

Modality	Model	$R@10 \uparrow$	$R@20 \uparrow$	$N@10 \uparrow$	$N@20 \uparrow$	$UCBias \downarrow$
Visual Modality	GUME	0.0536	0.0844	0.0291	0.0370	0.0564
	+D-ISOLATOR	<u>0.0549</u>	<u>0.0877</u>	<u>0.0301</u>	<u>0.0385</u>	<u>0.0553</u>
	+U-ISOLATOR	0.0552	0.0885	0.0302	0.0388	0.0529
Textual Modality	GUME	0.0675	0.1024	0.0362	0.0451	0.0717
	+D-ISOLATOR	0.0684	0.1034	<u>0.0366</u>	<u>0.0456</u>	0.0709
	+U-ISOLATOR	0.0684	0.1051	0.0372	0.0466	0.0781
Multi-modalities	GUME	0.0682	0.1041	<u>0.0368</u>	0.0460	0.0664
	+D-ISOLATOR	<u>0.0690</u>	<u>0.1062</u>	<u>0.0368</u>	<u>0.0464</u>	<u>0.0653</u>
	+U-ISOLATOR	0.0695	0.1063	0.0371	0.0465	0.0646

with the vanilla models. *D-ISOLATOR* completely removes the impact of User-side Content Similarity for all users, highlighting that **eliminating this impact can effectively reduce the User-side Content Bias**.

- *U-ISOLATOR* effectively enhances recommendation performance while controlling the User-side Content Bias in most cases. *U-ISOLATOR* improves accuracy and reduces the User-side Content Bias for all GCF-based MRSs. Despite a slight increase in User-side Content Bias for VBPR on the Baby and Clothing datasets, it achieves an average accuracy improvement of 7.10% on them. These highlights that **strategically using the impact of User-side Content Similarity is crucial**.

5.3 Impacts of Modality

To explore the effects of different modalities on ISOLATOR, we compare the results of GUME on the Baby dataset across different modalities because GUME performed the best among the backbones we used. We employ the same evaluation metrics in Sec. 5.1 and Sec. 5.2 to validate the accuracy and debiasing capability, respectively. From Table 3, we can observe that:

- Different modalities do not diminish the effectiveness of ISOLATOR. ISOLATOR consistently boosts recommendation performance (higher *Recall* and *NDCG*) while reducing User-side Content Bias (lower *UCBias*) in single-modal settings, mirroring the trends seen in multi-modal settings. This highlights its **robust generalization and broad adaptability across various modality environments**.
- Multi-modal content outperforms single-modal content (higher *Recall* and *NDCG*) by leveraging diverse modalities for more comprehensive information. This is also why we employ multi-modal content in this work.
- The textual modality provides more accurate recommendations than the visual modality due to higher explicitness and information density, but amplifies User-side Content Bias.
- User-side Content Bias exists across all modalities, even in single-modality settings, highlighting its **pervasive nature irrespective of the modality used**.

5.4 Impacts of Strategies

ISOLATOR employs three strategies: one training strategy *TS* and two inference strategies *D-ISOLATOR* and *U-ISOLATOR*. To assess

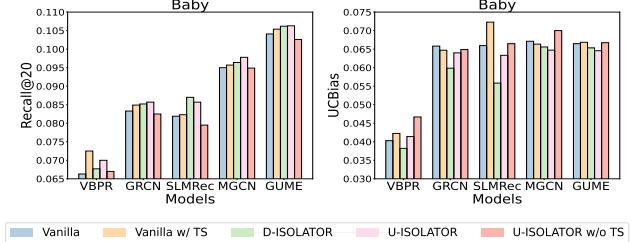


Figure 6: Performance and User-side Content Bias of Inference Strategies in ISOLATOR on the Baby Dataset: Higher Recall@20 and Lower UCBias indicate Better Results.

their effectiveness, we create these variants⁷: (1) *Vanilla*: no strategy in training or inference; (2) *Vanilla w/ TS*: TS in training, no strategy in inference; (3) *D-ISOLATOR*: TS in training, Debiasing intervention strategy in inference; (4) *U-ISOLATOR*: TS in training, User-specific intervention strategy in inference; (5) *U-ISOLATOR w/o TS*: no TS in training, User-specific intervention strategy in inference. We test these variants on the Baby dataset using the MRSs from Sec. 3.2.3 as backbones. We use the same evaluation metrics in Sec. 3.4. Higher *Recall@20* and lower *UCBias* indicate better results. From Fig. 6, we can observe that:

- **Combining training and inference strategies are the best.** *D-ISOLATOR* and *U-ISOLATOR* achieve higher *Recall@20* and lower *UCBias* than other variants across most backbones. Generally, ***D-ISOLATOR* keeps User-side Content Bias low, while *U-ISOLATOR* boosts accuracy**, enabling flexible strategy selection for accurate and unbiased recommendations.
- Training strategy boosts accuracy but also raises User-side Content Bias. All three variants using this strategy (i.e., *D-ISOLATOR*, *U-ISOLATOR*, *Vanilla w/ TS*) surpass the *vanilla* in *Recall@20*, yet *Vanilla w/ TS* gets a larger *UCBias*. This shows that training strategy can **model both the positive and negative impacts of User-side Content Similarity** but requires inference strategies to balance the trade-offs.
- Inference strategy can not be solely used. *U-ISOLATOR w/o TS* performs worse than others. Since the impact of User-side Content Similarity is not distinguished without training strategy, the inference strategy might overemphasize it, leading to poorer performance and increased bias.

6 Conclusion

In this paper, we uncover an **unexplored User-side Content Bias** in multi-modal recommender systems, and introduce an effective framework, ISOLATOR, to utilize multi-modal content more properly. Our empirical studies demonstrate the existence of User-side Content Bias, and our experiments validate the effectiveness of ISOLATOR. We provide insights into several key perspectives: (1) Utilizing User-side Content Similarity increases the exposure of items similar to the user’s past interactions, thereby improving the accuracy but amplifying User-side Content Bias. (2) Completely eliminating the impact of User-side Content Similarity can reduce User-side Content Bias, and strategically using the impact of User-side Content Similarity can enhance accuracy.

⁷*D-ISOLATOR w/o TS* equals *Vanilla*, as Debiasing intervention strategy doesn’t consider the impact of User-side Content Similarity.

Acknowledgments

Chen Lin is supported by the Natural Science Foundation of China (No.62372390,62432011).

References

- [1] Desheng Cai, Shengsheng Qian, Quan Fang, Jun Hu, and Changsheng Xu. 2022. Adaptive Anti-Bottleneck Multi-Modal Graph Learning Network for Personalized Micro-video Recommendation. In *MM '22: The 30th ACM International Conference on Multimedia, Lisboa, Portugal, October 10 - 14, 2022*. ACM, 581–590. doi:10.1145/3503161.3548420
- [2] Gaode Chen, Ruina Sun, Yuezhan Jiang, Jiangxia Cao, Qi Zhang, Jingjian Lin, Han Li, Kun Gai, and Xinghua Zhang. 2024. A Multi-modal Modeling Framework for Cold-start Short-video Recommendation. In *Proceedings of the 18th ACM Conference on Recommender Systems, RecSys 2024, Bari, Italy, October 14–18, 2024*. ACM, 391–400. doi:10.1145/3640457.3688098
- [3] Jingyuan Chen, Hanwang Zhang, Xiangnan He, Liqiang Nie, Wei Liu, and Tat-Seng Chua. 2017. Attentive Collaborative Filtering: Multimedia Recommendation with Item- and Component-Level Attention. In *Proceedings of the 40th International ACM SIGIR Conference on Research and Development in Information Retrieval, Shinjuku, Tokyo, Japan, August 7–11, 2017*. ACM, 335–344. doi:10.1145/3077136.3080797
- [4] Yewang Chen, Weiyao Ye, Guipeng Xv, Chen Lin, and Xiaomin Zhu. 2023. TCCM: Time and Content-Aware Causal Model for Unbiased News Recommendation. In *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management, CIKM 2023, Birmingham, United Kingdom, October 21–25, 2023*. ACM, 3778–3782. doi:10.1145/3583780.3615272
- [5] Djork-Arné Clevert, Thomas Unterthiner, and Sepp Hochreiter. 2016. Fast and Accurate Deep Network Learning by Exponential Linear Units (ELUs). In *4th International Conference on Learning Representations*.
- [6] Yashar Deldjoo, Vito Walter Anelli, Hamed Zamani, Alejandro Bellogín, and Tommaso Di Noia. 2021. A flexible framework for evaluating user and item fairness in recommender systems. *User Model. User Adapt. Interact.* 31, 3 (2021), 457–511. doi:10.1007/S11257-020-09285-1
- [7] Xavier Glorot and Yoshua Bengio. 2010. Understanding the difficulty of training deep feedforward neural networks. In *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics, AISTATS 2010, Chia Laguna Resort, Sardinia, Italy, May 13–15, 2010 (JMLR Proceedings, Vol. 9)*. JMLR.org, 249–256. <http://proceedings.mlr.press/v9/glorot10a.html>
- [8] Ruining He and Julian J. McAuley. 2016. BPVR: Visual Bayesian Personalized Ranking from Implicit Feedback. In *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence, February 12–17, 2016, Phoenix, Arizona, USA*. AAAI Press, 144–150. doi:10.1609/AAAI.V30I1.9973
- [9] Xiangnan He, Kuan Deng, Xiang Wang, Yan Li, Yong-Dong Zhang, and Meng Wang. 2020. LightGCN: Simplifying and Powering Graph Convolution Network for Recommendation. In *Proceedings of the 43rd International ACM SIGIR conference on research and development in Information Retrieval, SIGIR 2020, Virtual Event, China, July 25–30, 2020*. ACM, 639–648. <https://doi.org/10.1145/3397271.3401063>
- [10] Junjie Jia, Fen Wang, Huijuan Wang, and Shilong Liu. 2024. A Fairness Group Recommendation Algorithm Based On User Activity. *Int. J. Comput. Intell. Syst.* 17, 1 (2024), 204. doi:10.1007/S44196-024-00602-9
- [11] Tahsin Alamgir Khey, Mohamed Reda Bouadjenek, and Sunil Aryal. 2025. Unmasking Gender Bias in Recommendation Systems and Enhancing Category-Aware Fairness. In *Proceedings of the ACM on Web Conference 2025, WWW 2025, Sydney, NSW, Australia, 28 April 2025–2 May 2025*. ACM, 5127–5138. doi:10.1145/3696410.3714528
- [12] Diederik P. Kingma and Jimmy Ba. 2015. Adam: A Method for Stochastic Optimization. In *3rd International Conference on Learning Representations, ICLR 2015, San Diego, CA, USA, May 7–9, 2015, Conference Track Proceedings*. <http://arxiv.org/abs/1412.6980>
- [13] Dominik Kowald and Emanuel Lasic. 2022. Popularity Bias in Collaborative Filtering-Based Multimedia Recommender Systems. In *Advances in Bias and Fairness in Information Retrieval - Third International Workshop, BIAS 2022, Stavanger, Norway, April 10, 2022, Revised Selected Papers (Communications in Computer and Information Science, Vol. 1610)*. Springer, 1–11. doi:10.1007/978-3-031-09316-6_1
- [14] Kushal Kumar, Tarik Arici, Tal Neiman, Jinyu Yang, Shioulin Sam, Yi Xu, Hakan Ferhatosmanoglu, and Ismail B. Tutar. 2023. Unsupervised Multi-Modal Representation Learning for High Quality Retrieval of Similar Products at E-commerce Scale. In *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management, CIKM 2023, Birmingham, United Kingdom, October 21–25, 2023*. ACM, 4667–4673. doi:10.1145/3583780.3615504
- [15] Yunqi Li, Hanxiong Chen, Zuohui Fu, Yingqiang Ge, and Yongfeng Zhang. 2021. User-oriented Fairness in Recommendation. In *Proceedings of the Web Conference 2021 (Ljubljana, Slovenia) (WWW '21)*. Association for Computing Machinery, New York, NY, USA, 624–632. doi:10.1145/3442381.3449866
- [16] Yang Li, Qi'ao Zhao, Chen Lin, Jinsong Su, and Zhihui Zhang. 2024. Who To Align With: Feedback-Oriented Multi-Modal Alignment in Recommendation Systems. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2024, Washington DC, USA, July 14–18, 2024*. ACM, 667–676. doi:10.1145/3626772.36577701
- [17] Chen Lin, Xinyi Liu, Guipeng Xv, and Hui Li. 2021. Mitigating Sentiment Bias for Recommender Systems. In *SIGIR '21: The 44th International ACM SIGIR Conference on Research and Development in Information Retrieval, Virtual Event, Canada, July 11–15, 2021*. ACM, 31–40. doi:10.1145/3404835.3462943
- [18] Guojiao Lin, Zhen Meng, Dongjie Wang, Qingqing Long, Yuanchun Zhou, and Meng Xiao. 2024. GUME: Graphs and User Modalities Enhancement for Long-Tail Multimodal Recommendation. In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management, CIKM 2024, Boise, ID, USA, October 21–25, 2024*. ACM, 1400–1409. doi:10.1145/3627673.3679620
- [19] Chang Liu, Peng Hou, Anxiang Zeng, and Han Yu. 2024. Transformer-empowered multi-modal item embedding for enhanced image search in e-commerce. In *Proceedings of the Thirty-Eighth AAAI Conference on Artificial Intelligence and Thirty-Sixth Conference on Innovative Applications of Artificial Intelligence and Fourteenth Symposium on Educational Advances in Artificial Intelligence (AAAI'24/IAAI'24/EAAI'24)*. AAAI Press, Article 2588, 9 pages. doi:10.1609/aaai.v38i21.30311
- [20] Xinyi Liu, Wanxian Guo, Liyanuy Li, Hui Li, Chen Lin, Xubin Li, Si Chen, Jian Xu, Hongbo Deng, and Bo Zheng. 2022. Pretraining Representations of Multi-modal Multi-query E-commerce Search. In *KDD '22: The 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, Washington, DC, USA, August 14 – 18, 2022*. ACM, 3429–3437. doi:10.1145/3534678.3539200
- [21] Xiaohao Liu, Zhulin Tao, Jiahong Shao, Lifang Yang, and Xianglin Huang. 2022. EliMRec: Eliminating Single-modal Bias in Multimedia Recommendation. In *MM '22: The 30th ACM International Conference on Multimedia, Lisboa, Portugal, October 10 – 14, 2022*. ACM, 687–695. doi:10.1145/3503161.3548404
- [22] Daniele Malatesta, Giandomenico Cornacchia, Claudio Pomo, and Tommaso Di Noia. 2023. On Popularity Bias of Multimodal-aware Recommender Systems: A Modalities-driven Analysis. In *Proceedings of the 1st International Workshop on Deep Multimodal Learning for Information Retrieval, MMIR 2023, Ottawa ON, Canada, 2 November 2023*. ACM, 59–68. doi:10.1145/3606040.3617441
- [23] Jianmo Ni, Jiacheng Li, and Julian J. McAuley. 2019. Justifying Recommendations using Distantly-Labeled Reviews and Fine-Grained Aspects. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, EMNLP-IJCNLP 2019, Hong Kong, China, November 3–7, 2019*. Association for Computational Linguistics, 188–197. doi:10.18653/V1/D19-1018
- [24] Yongxin Ni, Yu Cheng, Xiangyan Liu, Junchen Fu, Youhua Li, Xiangnan He, Yongfeng Zhang, and Fajiu Yuan. 2023. A Content-Driven Micro-Video Recommendation Dataset at Scale. *arXiv preprint arXiv:2309.15379* (2023).
- [25] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Z. Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. 2019. PyTorch: An Imperative Style, High-Performance Deep Learning Library. In *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8–14, 2019, Vancouver, BC, Canada*. 8024–8035. <https://proceedings.neurips.cc/paper/2019/hash/bdbca288fe7f92f2bfa9f7012727740-Abstract.html>
- [26] Judea Pearl. 2009. *Causality*. Cambridge University Press.
- [27] YU Peng, Zhiyi Tan, Guanning Lu, and Bingkun Bao. 2023. Multi-View Graph Convolutional Network for Multimedia Recommendation. *Proceedings of the 31st ACM International Conference on Multimedia* (2023). <https://api.semanticscholar.org/CorpusID:260681274>
- [28] Steffen Rendle, Christoph Freudenthaler, Zeno Ganter, and Lars Schmidt-Thieme. 2009. BPR: Bayesian Personalized Ranking from Implicit Feedback. In *UAI 2009, Proceedings of the Twenty-Fifth Conference on Uncertainty in Artificial Intelligence, Montreal, QC, Canada, June 18–21, 2009*. AUAI Press, 452–461. https://www.auai.org/uai2009/papers/UAI2009_0139_4814db02b9f0b02bc7158819ebfa2c7.pdf
- [29] Ya Shang, Chen Gao, Jiansheng Chen, Depeng Jin, and Yong Li. 2024. Improving Item-side Fairness of Multimodal Recommendation via Modality Debiasing. In *Proceedings of the ACM on Web Conference 2024, WWW 2024, Singapore, May 13–17, 2024*. ACM, 4697–4705. doi:10.1145/3589334.3648156
- [30] Ya Shang, Chen Gao, Jiansheng Chen, Depeng Jin, Meng Wang, and Yong Li. 2023. Learning Fine-grained User Interests for Micro-video Recommendation. In *Proceedings of the 46th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2023, Taipei, Taiwan, July 23–27, 2023*. ACM, 433–442. <https://doi.org/10.1145/3539618.3591713>
- [31] Anne Shelley. 2012. Book review of Eli Pariser's *The filter bubble: What the Internet is hiding from you. First Monday* 17, 6 (2012). doi:10.5210/FM.V17I6.4100
- [32] Zhulin Tao, Xiaohao Liu, Yewei Xia, Xiang Wang, Lifang Yang, Xianglin Huang, and Tat-Seng Chua. 2023. Self-Supervised Learning for Multimedia Recommendation. *IEEE Trans. Multim.* 25 (2023), 5107–5116. doi:10.1109/TMM.2022.3187556
- [33] Wenjie Wang, Ful Feng, Xiangnan He, Xiang Wang, and Tat-Seng Chua. 2021. Deconfounded Recommendation for Alleviating Bias Amplification. In *KDD '21*

- The 27th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, Virtual Event, Singapore, August 14–18, 2021.* ACM, 1717–1725. doi:10.1145/3447548.3467249
- [34] Zongyi Wang, Yanyan Zou, Anyu Dai, Linfang Hou, Nan Qiao, Luobao Zou, Mian Ma, Zhuoye Ding, and Sulong Xu. 2023. An Industrial Framework for Personalized Serendipitous Recommendation in E-commerce. In *Proceedings of the 17th ACM Conference on Recommender Systems, RecSys 2023, Singapore, Singapore, September 18–22, 2023.* ACM, 1015–1018. <https://doi.org/10.1145/3604915.3610234>
- [35] Yinwei Wei, Xiang Wang, Liqiang Nie, Xiangnan He, and Tat-Seng Chua. 2020. Graph-Refined Convolutional Network for Multimedia Recommendation with Implicit Feedback. In *MM '20: The 28th ACM International Conference on Multimedia, Virtual Event / Seattle, WA, USA, October 12–16, 2020.* ACM, 3541–3549. doi:10.1145/3394171.3413556
- [36] Yinwei Wei, Xiang Wang, Liqiang Nie, Xiangnan He, Richang Hong, and Tat-Seng Chua. 2019. MMGCN: Multi-modal Graph Convolution Network for Personalized Recommendation of Micro-video. In *Proceedings of the 27th ACM International Conference on Multimedia, MM 2019, Nice, France, October 21–25, 2019.* ACM, 1437–1445. doi:10.1145/3343031.3351034
- [37] Guipeng Xv, Xinyu Li, Ruobing Xie, Chen Lin, Chong Liu, Feng Xia, Zhanhui Kang, and Leyu Lin. 2024. Improving Multi-modal Recommender Systems by Denoising and Aligning Multi-modal Content and User Feedback. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD 2024, Barcelona, Spain, August 25–29, 2024.* ACM, 3645–3656. doi:10.1145/3637528.3671703
- [38] Guipeng Xv, Chen Lin, Hui Li, Jinsong Su, Weiyao Ye, and Yewang Chen. 2022. Neutralizing Popularity Bias in Recommendation Models. In *SIGIR '22: The 45th International ACM SIGIR Conference on Research and Development in Information Retrieval, Madrid, Spain, July 11–15, 2022.* ACM, 2623–2628. doi:10.1145/3477495.3531907
- [39] Zixuan Yi, Xi Wang, Iadh Ounis, and Craig MacDonald. 2022. Multi-modal Graph Contrastive Learning for Micro-video Recommendation. In *SIGIR '22: The 45th International ACM SIGIR Conference on Research and Development in Information Retrieval, Madrid, Spain, July 11–15, 2022.* ACM, 1807–1811. doi:10.1145/3477495.3532027
- [40] Jinghao Zhang, Yanqiao Zhu, Qiang Liu, Shu Wu, Shuhui Wang, and Liang Wang. 2021. Mining Latent Structures for Multimedia Recommendation. In *MM '21: ACM Multimedia Conference, Virtual Event, China, October 20 – 24, 2021.* ACM, 3872–3880. doi:10.1145/3474085.3475259
- [41] Jinghao Zhang, Yanqiao Zhu, Qiang Liu, Mengqi Zhang, Shu Wu, and Liang Wang. 2023. Latent Structure Mining With Contrastive Modality Fusion for Multimedia Recommendation. *IEEE Trans. Knowl. Data Eng.* 35, 9 (2023), 9154–9167. doi:10.1109/TKDE.2022.32221949
- [42] Siqing Zhang, Yuchen Ding, Wei Tang, Wei Sun, Yong Liao, and Peng Yuan Zhou. 2025. Privacy-Preserving Orthogonal Aggregation for Guaranteeing Gender Fairness in Federated Recommendation. In *Proceedings of the Eighteenth ACM International Conference on Web Search and Data Mining, WSDM 2025, Hannover, Germany, March 10–14, 2025.* ACM, 327–335. doi:10.1145/3701551.3703513
- [43] Yang Zhang, Fulij Feng, Xiangnan He, Tianxin Wei, Chonggang Song, Guohui Ling, and Yongdong Zhang. 2021. Causal Intervention for Leveraging Popularity Bias in Recommendation. In *SIGIR '21: The 44th International ACM SIGIR Conference on Research and Development in Information Retrieval, Virtual Event, Canada, July 11–15, 2021.* ACM, 11–20. doi:10.1145/3404835.3462875
- [44] Xin Zhou. 2023. MMRec: Simplifying Multimodal Recommendation. In *ACM Multimedia Asia Workshops, MMAAsia 2023, Tainan, Taiwan, December 6–8, 2023.* ACM, 6:1–6:2. doi:10.1145/3611380.3628561
- [45] Xin Zhou and Zhiqi Shen. 2023. A Tale of Two Graphs: Freezing and Denoising Graph Structures for Multimodal Recommendation. In *Proceedings of the 31st ACM International Conference on Multimedia, MM 2023, Ottawa, ON, Canada, 29 October 2023 – 3 November 2023.* ACM, 935–943. doi:10.1145/3581783.3611943
- [46] Xin Zhou, Hongyu Zhou, Yong Liu, Zhiwei Zeng, Chunyan Miao, Pengwei Wang, Yuan You, and Feijun Jiang. 2023. Bootstrap Latent Representations for Multimodal Recommendation. In *Proceedings of the ACM Web Conference 2023, WWW 2023, Austin, TX, USA, 30 April 2023 – 4 May 2023.* ACM, 845–854. doi:10.1145/3543507.3583251